

# K-means Clustering

Rishabh Vaish

## K-Means Clustering

Let's consider coding a K-means algorithm. Perform the following:

- Load the `zip.train` (handwritten digit recognition) data from the `ElemStatLearn` package, and the goal is to identify clusters of digits based on only the pixels.

```
#load data  
library(ElemStatLearn)  
data("zip.train")  
#remove the first column  
x <- zip.train[, 2:ncol(zip.train)]
```

- Writing your own code of k-means that iterates between two steps, and stop when the cluster membership does not change.
  - updating the cluster means given the cluster membership
  - updating the cluster membership based on cluster means

```
# k means clustering function takes x and k as input and gives  
# final clusters for each x_i as output  
k_means <- function(x, k) {  
  # assigning random clusters initially  
  clustered <-  
    as.matrix(cbind(x, cluster = floor(runif(  
      nrow(x), min = 1, max = 6  
    ))))  
  
  new_clusters = rep(0, nrow(clustered))  
  iter = 0  
  
  # while loop to assign and update clusters till  
  # clusters dont change or maximum iterations is reached  
  while (iter < 100) {  
    #calculate cluster means  
    cluster_means <-  
      as.matrix(aggregate(clustered, list(clustered[, 257]), mean))  
    #remove group column  
    cluster_means <- cluster_means[,-1]  
    # assigning new clusters to each row on the basis of cluster means  
    for (i in 1:nrow(clustered)) {  
      distances <- sqrt(rowSums(sweep(cluster_means, 2, clustered[i, ]) ^ 2))  
      new_clusters[i] <- which.min(distances)
```

```

}
#checking if the clusters change, else update clusters
if (all(new_clusters == clustered[, ncol(clustered)])) {
  break
} else{
  clustered[, ncol(clustered)] = new_clusters
}
#take count of iterations
iter = iter + 1
}
#return new clusters
return(new_clusters)
}

```

- Perform your algorithm with one random initialization with  $k = 5$ 
  - For this question, compare your cluster membership to the true digits. What are the most prevalent digits in each of your clusters?

```

#one iteration
k = 5
start_time = Sys.time()
clust <- k_means(x, k)
end_time = Sys.time()
#time taken for one iteration
end_time - start_time

```

## Time difference of 10.4003 secs

```

#most prevalent digits per cluster
results <-
  as.data.frame(cbind(cluster = clust, digit = zip.train[, 1]))
#clusters vs digits analysis
table(results)

```

```

##      digit
## cluster  0   1   2   3   4   5   6   7   8   9
##      1  25   0 108 622   4 381   0   2 347  11
##      2 252   1 531  19 203 117 569   9  52  12
##      3 915   0  31   2   4  42  83   1   8   3
##      4   0 1004   7   2  55   0  12   7  30  25
##      5   2   0  54  13 386  16   0 626 105 593

```

- Perform your algorithm with 10 independent initiations with  $k = 5$  and record the best
  - For this question, plot your clustering results on a two-dimensional plot, where the two axis are the first two principle components of your data

```

#for 10 independent initiations
clust <- matrix(0, nrow = nrow(x), ncol = 10)
for (i in 1:10) {
  clust[, i] <- k_means(x, k)
}

```

```

}

#calculate between cluster distance for each of the iterations
# to choose best model
between_clust_dist <- rep(0, 10)
for (i in 1:10) {
  #calculate cluster means
  data <- as.matrix(cbind(x, cluster = clust[, i]))
  cluster_mean <- as.matrix(aggregate(data, list(data[, 257]), mean))
  sum = 0
  count = 1
  #calculate average between cluster distance for each iteration
  for (m in 1:k) {
    for (n in m:k) {
      sum = sum + sqrt(sum((cluster_mean[m, 1:256] - cluster_mean[n, 1:256]) ^
                           2))
      count = count + 1
    }
  }
  between_clust_dist[i] <- sum / count
}

which.max(between_clust_dist)

```

```
## [1] 3
```

```

#pca
zip_pc = prcomp(x)

#plot on pca components
library(ggplot2)
library(reshape2)
ggplot(data = data.frame(zip_pc$x), aes(x = PC1, y = PC2)) +
  geom_point(color = rainbow_hcl(5)[clust[, which.max(between_clust_dist)]], size = 1)

```



- Compare the clustering results from the above two questions with the built-in `kmeans()` function in R. Use tables/figures to demonstrate your results and comment on your findings.

```

set.seed(1)

# k mean clustering
kmean <- kmeans(zip.train[, 2:ncol(zip.train)], centers = 5, nstart = 10)

#most prevalant didgit per cluster for r kmeans models
result_r <-
  as.data.frame(cbind(cluster = kmean$cluster, digit = zip.train[, 1]))
#clusters vs digits analysis
result_r <- table(result_r)
result_r

```

```

##      digit
## cluster 0    1    2    3    4    5    6    7    8    9
##      1   0 1003    6    2   59    1   13    7   23   23
##      2    1    0   48   21  484   32    3  629  190  609
##      3  881    0   22    2    4   36   68    1    5    3
##      4   28    1   93  615    4  383    2    1  262    4
##      5  284    1  562   18  101  104  578    7   62    5

```

```

#most prevalant didgit per cluster for best of 10 models
result <-
  as.data.frame(cbind(cluster = clust[, which.max(between_clust_dist)], digit = zip.train[, 1]))
#clusters vs digits analysis
table(result)

```

```

##      digit
## cluster 0    1    2    3    4    5    6    7    8    9
##      1  416    0  108   25   28  198  508    1   20    3
##      2    5    0   74   13  494   38    7  631   71  599
##      3    1 1003   55    2   83    2   37    7   29   33
##      4  732    0   22    6   11   50   66    0    7    3
##      5   40    2  472  612   36  268   46    6  415    6

```

According to R Kmeans following are the digits in each clusters - *cluster1* - {1} *cluster2* - {4,7,9} *cluster3* - {0} *cluster4* - {3,5,8} \**cluster5* - {2,6}

According to 10 iteration with my function following are the digits in each clusters of the best model- *cluster1* - {1} *cluster5* - {4,7,9} *cluster2* - {0} *cluster3* - {2,3,8} \**cluster4* - {5,6}

According to 1 iteration with my function following are the digits in each clusters - *cluster2* - {1} *cluster4* - {4,7,9} *cluster3* - {0} *cluster5* - {3,5,6,8} \**cluster1* - {2}

The best of 10 models using my function works nearly as good as the kmeans R function. Just 2 and 5 are the digits swapped. The digits {4,7,9} are similar and are rightly clustered together. {0} and {1} is not clustered with any other digit, rightly so. In my algo {5,6} are clustered together whereas {2,6} are clustered together in R Kmeans. Results are mostly similar with an exception of 2 and 5. The clustering seems correct and justifiable.